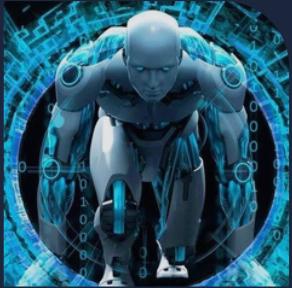


Synthetic Segmentation Training Data Generator

| Project Overview |

AI-Powered Innovation

+ Content:



ABOUT US
INTRODUCTION
DEEP DIVE INTO INSIGHTS
OVERVIEW OF LUNAPIX
THE UNIQUE APPROACH
EXPERIENCE LUNAPIX
THE FUTURE OF LUNAPIX
CONCLUSION





+ Introduction

PROJECT TITLE: SYNTHETIC SEGMENTATION TRAINING DATA GENERATOR

PROBLEM:

- TRAINING DATA IS HARD TO GET BY AND EXPENSIVE.
- DEVELOP A MODEL TO GENERATE HIGH-QUALITY SYNTHETIC DATA SETS

HOW:

- LEVERAGED THE USE OF GAN MODEL
- USING A DISCRIMINATOR
- WE ACHIEVE HIGH QUALITY ARTIFICIAL IMAGES WITH REALISTIC QUALITIES



+ BRIEF PROCESS OVERVIEW

STEP 1:

- RESEARCHING AND UNDERSTANDING POTENTIAL SOLUTIONS

STEP 2:

- DATA COLLECTION AND PREPARATION

STEP 3:

- MODEL DESIGN AND SETUP

STEP 4:

- TRAINING THE GAN

STEP 5:

- EVALUATION AND FINE-TUNING

STEP 6:

- DEPLOYMENT AND APPLICATION



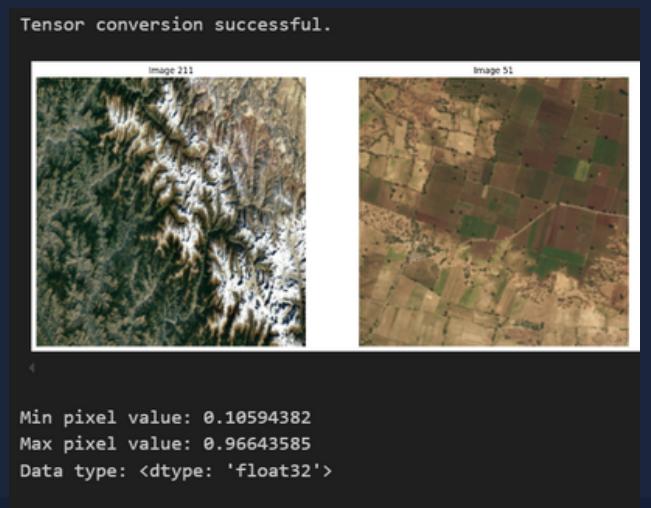
+ A CLOSER LOOK - RAW IMAGES

IMAGES HAD TO BE RESIZED, NORMALIZING PIXEL VALUES, AND ENSURING THE OVERLL AUGMENTATION OF THE DATASET IS SUITABLE .



TENSOR CONVERSION SUCCESSFUL:

- USING TENSORFLOW
- THE MINIMUM & MAXIMUM PIXEL VALUES
- THE TENSOR DATA TYPE - IS FLOAT32 FORMAT FOR NEURAL NETWORK INPUTS,



+ A CLOSER LOOK - TENSORFLOW CONVERSION

STEPS THE CODE FOLLOWS:

1. CONVERT AND SCALE IMAGES

2. RESIZED IMAGES

3. COLOUR CHANNELS OF THE IMAGES

4. TENSOR CONVERSION WITH RANDOM DISPLAY

5. PIXEL VALUE AND DATA TYPE CHECK

```
1 # Preprocessing images - adjusting size and format to make them ready for further processing
2
3 def preprocess_image(image_array, target_size=(256, 256)):
4     # Convert the image to float32 to reduce memory consumption
5     image_array = image_array.astype(np.float32) / 255.0
6
7     # Resize image
8     image_resized = resize(image_array, target_size, anti_aliasing=True)
9
10    # Standardize color channels: ensure all images are in RGB format
11    if image_resized.ndim == 2: # If the image is grayscale, convert it to RGB
12        image_resized = np.stack([image_resized] * 3, axis=-1)
13    elif image_resized.shape[-1] == 4: # If the image has an alpha channel, remove it
14        image_resized = image_resized[:, :, :-1]
15
16    return image_resized
17
18 # Process list of images individually, applying transformations to each image for consistent size and colour format
19 # Managing memory usage effectively & prevent crashes due to memory overflow
20 images = []
21 for img in image_list:
22     try:
23         processed_image = preprocess_image(img)
24         images.append(processed_image)
25     except MemoryError as e:
26         print(f"Memory error while processing the image: {e}")
27     # Handle the memory error, (e.g., by skipping the image or freeing up resources)
28     break
29
30 # Convert list of processed images into a TensorFlow tensor
31 # Displays five random images from this tensor; if no images are processed - print a statement to inform
32 if images:
33     images_tensor = tf.convert_to_tensor(images, dtype=tf.float32)
34     print("Tensor conversion successful.")
35
36 # Display 5 random images
37 def display_random_images(images, num_images=5):
38     plt.figure(figsize=(40, 20))
39     indices = np.random.choice(images.shape[0], num_images, replace=False)
40     for i, index in enumerate(indices):
41         ax = plt.subplot(1, num_images, i + 1)
42         plt.imshow(images[index])
43         plt.title(f"Image {index+1}")
44         plt.axis("off")
45     plt.show()
46
47 display_random_images(images)
48 else:
49     print("No images that were processed.")
50     print("Tensor conversion successful.")
51
52 # Verify the minimum and maximum pixel values and the data type
53 # Helps ensure that the image data is correctly formatted and scaled for processing tasks
54 print("Min pixel value:", np.min(images_tensor[0]))
55 print("Max pixel value:", np.max(images_tensor[0]))
56
57
58 # Ensure that the dtype is float32
59 print("Data type:", images_tensor.dtype)
```

CREATING A DATAFRAME FOR SEGMENTING IMAGES

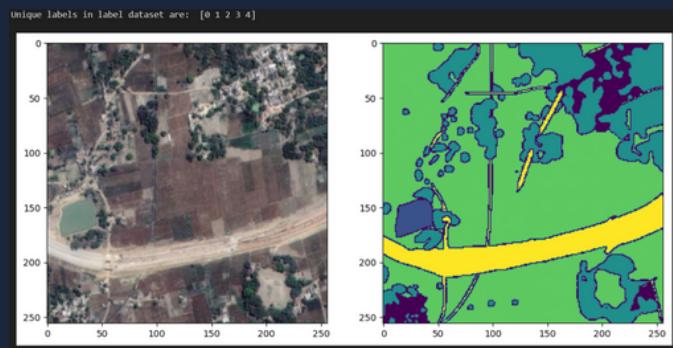
- CREATING A DATAFRAME FOR SEGMENTING IMAGES.
- LABELING THE IMAGES
- DISPLAYING RGB VALUES OF THE DIFFERENT IMAGE LABELS

```
1 import pandas as pd
2 from io import StringIO
3
4     # Specify the bucket name and file name
5 bucket_name = '2307-07-synthetic-segmentation-a'
6 file_key = 'class_dict_seg.csv'
7
8 #Initialize S3 client
9 s3_client = boto3.client(
10     service_name='s3',
11     region_name='eu-west-1',
12     aws_access_key_id=os.getenv("aws_access_key_id"),
13     aws_secret_access_key=os.getenv("aws_secret_access_key"))
14
15 def read_csv_from_s3(bucket_name, file_key):
16
17     # Read CSV file from S3
18     obj = s3_client.get_object(Bucket=bucket_name, Key=file_key)
19     data = obj['Body'].read().decode('utf-8')
20
21     # Convert CSV data to Pandas DataFrame
22     mask_labels = pd.read_csv(StringIO(data))
23
24     return mask_labels
25
26 mask_labels = read_csv_from_s3(bucket_name, file_key)
27 print(mask_labels.head()) # Display the first few rows of the DataFrame
```

	name	r	g	b
0	urban	0	255	255
1	water	0	0	255
2	forest	0	255	0
3	land	255	255	0
4	road	255	0	255

CONVERTING RGB IMAGES & SEGMENTING

- CONVERTING RGB IMAGES
- CLASS LABELING OF THE IMAGES
- DISPLAYING RGB VALUES OF IMAGE LABELS



```
# converting 3 values (RGB) to 1 label values...
def rgb_to_labels(img, mask_labels):
    label_seg = np.zeros(img.shape,dtype=np.uint8)
    for i in range(mask_labels.shape[0]):
        label_seg[np.all(img == list(mask_labels.iloc[i, [1,2,3]]), axis=-1)] = i
    label_seg = label_seg[:, :, 0] #Just take the first channel, no need for all 3 channels...
    return label_seg

labels = []
for i in range(mask_dataset.shape[0]):
    label = rgb_to_labels(mask_dataset[i], mask_labels) # calling rgb to labels for each images...
    labels.append(label)

labels = np.array(labels)
labels = np.expand_dims(labels, axis=3) # adding channel dim...
```



LABEL STUDIO

DEMO IMAGE
FROM LABEL
STUDIO

Original Image

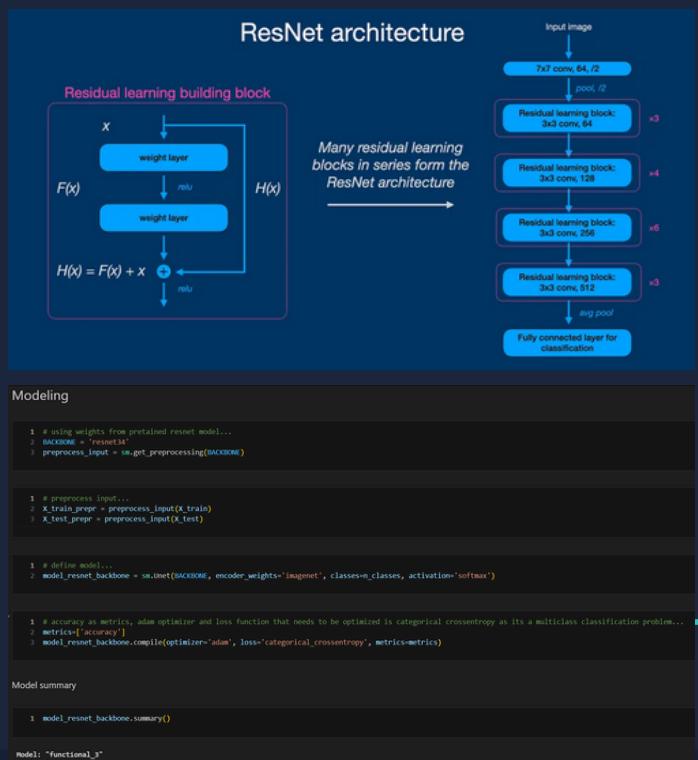


Binarized Segmented Image



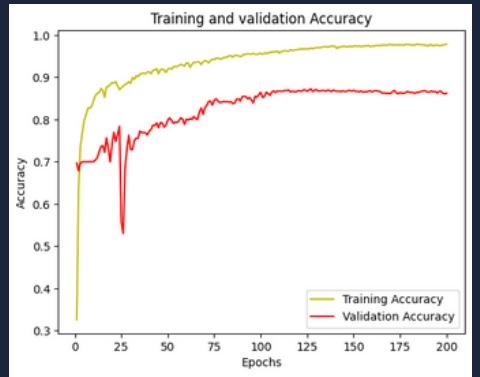
TRAINING THE SEGMENTATION MODEL

- RESNET MODEL
- DEFINING THE MODEL
- SUMMARY OF THE MODEL ARCHITECTURE

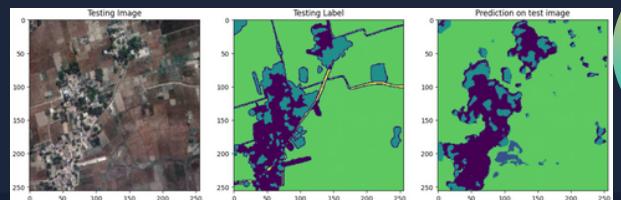


TRAINING RESULTS AFTER 200 EPOCHS

- ACCURACY - TRAINING
- CODE INITIATION AND SETTING PARAMETERS
- TRAINING ACCURACY - 0.95
- VALIDATION ACCURACY- 0.86
- LOSS FUNCTION - 0.12



```
# training the model for 80 epochs with batch size of 16...
history = model_resnet_backbone.fit(X_train_prep,
y_train,
batch_size=16,
epochs=200,
verbose=1,
validation_data=(X_test_prep, y_test))
```



BUILDING A GENERATOR MODEL

THE ARTIST THE GENERATOR



```
1 from tensorflow.keras.layers import Conv2D, Conv2DTranspose, BatchNormalization, LeakyReLU, Activation, Concatenate
2 from tensorflow.keras import Input, Model
3
4 def build_generator():
5     inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
6
7     # Encoder
8     e1 = Conv2D(64, (4, 4), strides=(2, 2), padding='same')(inputs)
9     e1 = LeakyReLU(alpha=0.2)(e1)
10
11    e2 = Conv2D(128, (4, 4), strides=(2, 2), padding='same')(e1)
12    e2 = BatchNormalization()(e2)
13    e2 = LeakyReLU(alpha=0.2)(e2)
14
15    e3 = Conv2D(256, (4, 4), strides=(2, 2), padding='same')(e2)
16    e3 = BatchNormalization()(e3)
17    e3 = LeakyReLU(alpha=0.2)(e3)
18
19    # Decoder
20    d1 = Conv2DTranspose(128, (4, 4), strides=(2, 2), padding='same')(e3)
21    d1 = BatchNormalization()(d1)
22    d1 = Activation('relu')(d1)
23
24    d2 = Conv2DTranspose(64, (4, 4), strides=(2, 2), padding='same')(d1)
25    d2 = BatchNormalization()(d2)
26    d2 = Activation('relu')(d2)
27
28    outputs = Conv2DTranspose(IMG_CHANNELS, (4, 4), strides=(2, 2), padding='same')(d2)
29    outputs = Activation('tanh')(outputs)
30
31    return Model(inputs, outputs)
```

BUILDING A DISCRIMINATOR

THE DISCRIMINATOR ARCHITECTURE

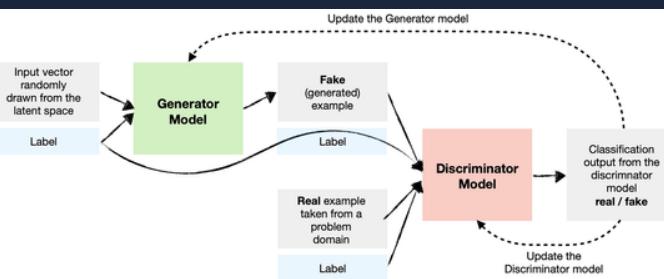


```
# Discriminator
def build_discriminator():
    inputs = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
    targets = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
    combined = Concatenate()([inputs, targets])
    # Add layers for the discriminator
    down1 = Conv2D(64, kernel_size=4, strides=2, padding='same')(combined)
    down1 = LeakyReLU(alpha=0.2)(down1)
    down2 = Conv2D(128, kernel_size=4, strides=2, padding='same')(down1)
    down2 = BatchNormalization()(down2)
    down2 = LeakyReLU(alpha=0.2)(down2)
    down3 = Conv2D(256, kernel_size=4, strides=2, padding='same')(down2)
    down3 = BatchNormalization()(down3)
    down3 = LeakyReLU(alpha=0.2)(down3)
    down4 = Conv2D(512, kernel_size=4, strides=2, padding='same')(down3)
    down4 = BatchNormalization()(down4)
    down4 = LeakyReLU(alpha=0.2)(down4)
    outputs = Conv2D(1, kernel_size=4, strides=1, padding='same', activation='sigmoid')(down4)
    return Model([inputs, targets], outputs)
```



COMBINING THE TWO

SYNERGISED PROCESS



```
1  from tensorflow.keras.optimizers import Adam
2  from tensorflow.keras.models import Model
3
4  # Compile discriminator
5  discriminator.compile(loss='binary_crossentropy', optimizer=Adam(2e-4, beta_1=0.5), metrics=['accuracy'])
6
7  # For the combined model, only train the generator
8  discriminator.trainable = False
9
10 # Input images and their conditioning images
11 input_img = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS))
12 gen_output = generator(input_img)
13
14 # Discriminator determines validity of translated images / condition pairs
15 dis_output = discriminator([input_img, gen_output])
16
17 # Combined model (stacked generator and discriminator)
18 combined = Model(inputs=input_img, outputs=[dis_output, gen_output])
19 combined.compile(loss=['binary_crossentropy', 'mae'], optimizer=Adam(2e-4, beta_1=0.5))
```

TRAINING THE MODEL TO GENERATE SYNTHETIC IMAGES

MODEL IS NOW TRAINED TO DISPLAY
REAL, FAKE AND MASK IMAGES

```
1/1      0s 288ms/step
[Epoch 301/311] [D loss: 5.780974388122559] [G loss: 5.7764739990234375] time: 6.717052221298218
1/1      0s 292ms/step
[Epoch 302/311] [D loss: 5.7811994552612305] [G loss: 5.776702880859375] time: 5.655746698379517
1/1      0s 294ms/step
[Epoch 303/311] [D loss: 5.777637481689453] [G loss: 5.773155689239502] time: 5.764461994171143
1/1      0s 291ms/step
[Epoch 304/311] [D loss: 5.778769493103027] [G loss: 5.774311542510986] time: 7.130734205245972
1/1      0s 290ms/step
[Epoch 305/311] [D loss: 5.779085159301758] [G loss: 5.774621963500977] time: 4.683357000350952
1/1      0s 270ms/step
[Epoch 306/311] [D loss: 5.780157566070557] [G loss: 5.775706768035889] time: 4.686371564865112
1/1      0s 261ms/step
[Epoch 307/311] [D loss: 5.781589588056641] [G loss: 5.777185440063477] time: 6.669575214385986
1/1      0s 315ms/step
[Epoch 308/311] [D loss: 5.782863616943359] [G loss: 5.778444766998291] time: 5.644251823425293
1/1      0s 291ms/step
[Epoch 309/311] [D loss: 5.783916473388672] [G loss: 5.779517650604248] time: 3.76678466796875
```

```
5  def train(epochs, batch_size, save_interval):
6      real = np.ones((batch_size, 16, 16, 1))
7      fake = np.zeros((batch_size, 16, 16, 1))
8
9      for epoch in range(epochs):
10         start_time = time.time()
11
12         # Train discriminator with a batch of real images
13         idx = np.random.randint(0, image_dataset.shape[0], batch_size)
14         real_imgs = image_dataset[idx]
15         masks = mask_dataset[idx]
16
17         fake_imgs = generator.predict(real_imgs)
18
19         d_loss_real = discriminator.train_on_batch([real_imgs, masks], real)
20         d_loss_fake = discriminator.train_on_batch([real_imgs, fake_imgs], fake)
21         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
22
23         # Train generator
24         g_loss = combined.train_on_batch(real_imgs, [real, masks])
25
26         elapsed_time = time.time() - start_time
27
28         # Print the progress
29         print(f"[epoch {epoch}]/({epochs}) [D loss: {d_loss[0]}] [G loss: {g_loss[0]}] time: {elapsed_time}")
30
31         # If at save interval, save generated image samples
32         if epoch % save_interval == 0:
33             save_model(epoch)
34             sample_images(epoch)
35
36     def sample_images(epoch):
37         r, c = 3, 3
38         idx = np.random.randint(0, image_dataset.shape[0], c)
39         real_imgs = image_dataset[idx]
40         masks = mask_dataset[idx]
41         fake_imgs = generator.predict(real_imgs)
```



+ DISPLAYING THE RESULTS

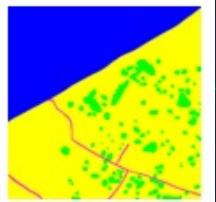
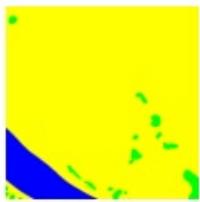
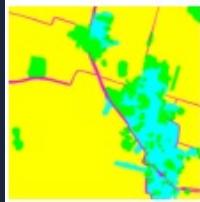
THE ORIGINAL



THE PREDICTION



THE CONDITION
(SEGMENTED)





THANK YOU